

Vi Le_Project 2

December 3, 2024

1 Project 2 - Airline Customer Satisfaction

1.1 Vi Le | DAT 402: Machine Learning- Data Science, Fall 2024

Customer satisfaction is a cornerstone of success for businesses in industries such as retail, telecommunications, and airlines. Satisfied customers are more likely to become loyal advocates, reflecting a brand's quality and trustworthiness.

With over four years of experience in retail and customer service, I have seen firsthand how critical it is to meet and exceed customer expectations. Inspired by this, I chose to analyze and predict customer satisfaction levels in the airline industry for this project. By leveraging data insights, my goal is to better understand the factors that drive satisfaction and loyalty, as well as to build a model that could classify happy and unhappy customers based on different factors.

Credits: - The following data set is from TJ Klein (link: <https://www.kaggle.com/datasets/teejmahal20/airline-passenger-satisfaction>) - The data set provided is modified from John D. (link: <https://www.kaggle.com/datasets/johnddddd/customer-satisfaction>)

In this project, I will use 3 different classification methods on the same data set, and compare their performances: - Naive Bayes - Logistic Regression - Decision Tree

First thing first, I will import all necessary libraries and packages.

```
[9]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix
from sklearn.naive_bayes import MultinomialNB, CategoricalNB, GaussianNB
from sklearn.metrics import confusion_matrix, accuracy_score
```

Then, I will load the data set, and see a first few rows of it.

```
[11]: df = pd.read_csv("airline.csv")
df = pd.DataFrame(df)
df.head()
```

```
[11]: Unnamed: 0      id  Gender  Customer Type  Age  Type of Travel \
0          0    70172   Male    Loyal Customer   13  Personal Travel
```

1	1	5047	Male	disloyal Customer	25	Business travel
2	2	110028	Female	Loyal Customer	26	Business travel
3	3	24026	Female	Loyal Customer	25	Business travel
4	4	119299	Male	Loyal Customer	61	Business travel

	Class	Flight Distance	Inflight wifi service	\
0	Eco Plus	460	3	
1	Business	235	3	
2	Business	1142	2	
3	Business	562	2	
4	Business	214	3	

	Departure/Arrival time convenient	...	Inflight entertainment	\
0	4	...	5	
1	2	...	1	
2	2	...	5	
3	5	...	2	
4	3	...	3	

	On-board service	Leg room service	Baggage handling	Checkin service	\
0	4	3	4	4	
1	1	5	3	1	
2	4	3	4	4	
3	2	5	3	1	
4	3	4	4	3	

	Inflight service	Cleanliness	Departure Delay in Minutes	\
0	5	5	25	
1	4	1	1	
2	4	5	0	
3	4	2	11	
4	3	3	0	

	Arrival Delay in Minutes	satisfaction
0	18.0	neutral or dissatisfied
1	6.0	neutral or dissatisfied
2	0.0	satisfied
3	9.0	neutral or dissatisfied
4	0.0	satisfied

[5 rows x 25 columns]

```
[12]: print(df.dtypes)
```

```
Unnamed: 0          int64
id                int64
Gender             object
Customer Type      object
```

Age	int64
Type of Travel	object
Class	object
Flight Distance	int64
Inflight wifi service	int64
Departure/Arrival time convenient	int64
Ease of Online booking	int64
Gate location	int64
Food and drink	int64
Online boarding	int64
Seat comfort	int64
Inflight entertainment	int64
On-board service	int64
Leg room service	int64
Baggage handling	int64
Checkin service	int64
Inflight service	int64
Cleanliness	int64
Departure Delay in Minutes	int64
Arrival Delay in Minutes	float64
satisfaction	object
dtype:	object

Explanation of column names: - Gender: Gender of the passengers (Female, Male) - Customer Type: The customer type (Loyal customer, disloyal customer) - Age: The actual age of the passengers - Type of Travel: Purpose of the flight of the passengers (Personal Travel, Business Travel) - Class: Travel class in the plane of the passengers (Business, Eco, Eco Plus) - Flight distance: The flight distance of this journey - Inflight wifi service: Satisfaction level of the inflight wifi service (0:Not Applicable;1-5) - Departure/Arrival time convenient: Satisfaction level of Departure/Arrival time convenient - Ease of Online booking: Satisfaction level of online booking - Gate location: Satisfaction level of Gate location - Food and drink: Satisfaction level of Food and drink - Online boarding: Satisfaction level of online boarding - Seat comfort: Satisfaction level of Seat comfort - Inflight entertainment: Satisfaction level of inflight entertainment - On-board service: Satisfaction level of On-board service - Leg room service: Satisfaction level of Leg room service - Baggage handling: Satisfaction level of baggage handling - Check-in service: Satisfaction level of Check-in service - Inflight service: Satisfaction level of inflight service - Cleanliness: Satisfaction level of Cleanliness - Departure Delay in Minutes: Minutes delayed when departure - Arrival Delay in Minutes: Minutes delayed when Arrival - Satisfaction: Airline satisfaction level(Satisfaction, neutral or dissatisfaction)

1.1.1 Step 2: Data Preprocessing

```
[19]: # Check for NA values
      print(df.isna().sum())
```

```
Unnamed: 0          0
id                0
```

```

Gender                                0
Customer Type                         0
Age                                   0
Type of Travel                        0
Class                                 0
Flight Distance                       0
Inflight wifi service                 0
Departure/Arrival time convenient    0
Ease of Online booking                0
Gate location                         0
Food and drink                        0
Online boarding                       0
Seat comfort                          0
Inflight entertainment                0
On-board service                      0
Leg room service                      0
Baggage handling                      0
Checkin service                       0
Inflight service                      0
Cleanliness                           0
Departure Delay in Minutes            0
Arrival Delay in Minutes              310
satisfaction                          0
dtype: int64

```

```
[21]: df = df.dropna()
```

```
[23]: # Check statistical aspect of data
df.describe()
```

```
[23]:
```

	Unnamed: 0	id	Age	Flight Distance \
count	103594.000000	103594.000000	103594.000000	103594.000000
mean	51950.102274	64942.428625	39.380466	1189.325202
std	29997.914016	37460.816597	15.113125	997.297235
min	0.000000	1.000000	7.000000	31.000000
25%	25960.250000	32562.250000	27.000000	414.000000
50%	51955.500000	64890.000000	40.000000	842.000000
75%	77924.750000	97370.500000	51.000000	1743.000000
max	103903.000000	129880.000000	85.000000	4983.000000

	Inflight wifi service	Departure/Arrival time convenient \
count	103594.000000	103594.000000
mean	2.729753	3.060081
std	1.327866	1.525233
min	0.000000	0.000000
25%	2.000000	2.000000
50%	3.000000	3.000000

75%	4.000000	4.000000
max	5.000000	5.000000

	Ease of Online booking	Gate location	Food and drink	Online boarding \
count	103594.000000	103594.000000	103594.000000	103594.000000
mean	2.756984	2.977026	3.202126	3.250497
std	1.398934	1.277723	1.329401	1.349433
min	0.000000	0.000000	0.000000	0.000000
25%	2.000000	2.000000	2.000000	2.000000
50%	3.000000	3.000000	3.000000	3.000000
75%	4.000000	4.000000	4.000000	4.000000
max	5.000000	5.000000	5.000000	5.000000

	Seat comfort	Inflight entertainment	On-board service \
count	103594.000000	103594.000000	103594.000000
mean	3.439765	3.358341	3.382609
std	1.318896	1.333030	1.288284
min	0.000000	0.000000	0.000000
25%	2.000000	2.000000	2.000000
50%	4.000000	4.000000	4.000000
75%	5.000000	4.000000	4.000000
max	5.000000	5.000000	5.000000

	Leg room service	Baggage handling	Checkin service	Inflight service \
count	103594.000000	103594.000000	103594.000000	103594.000000
mean	3.351401	3.631687	3.304323	3.640761
std	1.315409	1.181051	1.265396	1.175603
min	0.000000	1.000000	0.000000	0.000000
25%	2.000000	3.000000	3.000000	3.000000
50%	4.000000	4.000000	3.000000	4.000000
75%	4.000000	5.000000	4.000000	5.000000
max	5.000000	5.000000	5.000000	5.000000

	Cleanliness	Departure Delay in Minutes	Arrival Delay in Minutes
count	103594.000000	103594.000000	103594.000000
mean	3.286397	14.747939	15.178678
std	1.312194	38.116737	38.698682
min	0.000000	0.000000	0.000000
25%	2.000000	0.000000	0.000000
50%	3.000000	0.000000	0.000000
75%	4.000000	12.000000	13.000000
max	5.000000	1592.000000	1584.000000

“Customer Type”, “Type of Travle”, “Class”, “Gender”, and “satisfaction” are categorical types and contain text values. For best use of data for machine learning models, I will turn those values into numerical type using **OrdinalEncoder()**.

```
[28]: from sklearn.preprocessing import OrdinalEncoder
ordinal = OrdinalEncoder()
df['Customer Type'] = ordinal.fit_transform(df[['Customer Type']])
df['Type of Travel'] = ordinal.fit_transform(df[['Type of Travel']])
df['Class'] = ordinal.fit_transform(df[['Class']])
df['satisfaction'] = ordinal.fit_transform(df[['satisfaction']])
df['Gender'] = ordinal.fit_transform(df[['Gender']])
df.head()
```

```
[28]: Unnamed: 0      id  Gender  Customer Type  Age  Type of Travel  Class  \
0          0    70172     1.0           0.0   13           1.0    2.0
1          1     5047     1.0           1.0   25           0.0    0.0
2          2   110028     0.0           0.0   26           0.0    0.0
3          3    24026     0.0           0.0   25           0.0    0.0
4          4   119299     1.0           0.0   61           0.0    0.0

      Flight Distance  Inflight wifi service  Departure/Arrival time convenient  \
0                460                        3                               4
1                235                        3                               2
2               1142                        2                               2
3                562                        2                               5
4                214                        3                               3

...  Inflight entertainment  On-board service  Leg room service  \
0 ...                      5                  4                3
1 ...                      1                  1                5
2 ...                      5                  4                3
3 ...                      2                  2                5
4 ...                      3                  3                4

      Baggage handling  Checkin service  Inflight service  Cleanliness  \
0                4                4                5                5
1                3                1                4                1
2                4                4                4                5
3                3                1                4                2
4                4                3                3                3

      Departure Delay in Minutes  Arrival Delay in Minutes  satisfaction
0                25                18.0                0.0
1                 1                 6.0                0.0
2                 0                 0.0                1.0
3                11                 9.0                0.0
4                 0                 0.0                1.0
```

[5 rows x 25 columns]

1.1.2 Step 3: Visualize data for insights

```
[32]: corr_matrix = df.corr()  
      corr_matrix["satisfaction"].sort_values(ascending=False)
```

```
[32]: satisfaction          1.000000  
      Online boarding      0.503447  
      Inflight entertainment 0.398203  
      Seat comfort         0.349112  
      On-board service     0.322450  
      Leg room service     0.313182  
      Cleanliness          0.305050  
      Flight Distance      0.298915  
      Inflight wifi service 0.284163  
      Baggage handling     0.247819  
      Inflight service     0.244852  
      Checkin service      0.235914  
      Food and drink       0.209659  
      Ease of Online booking 0.171507  
      Age                 0.137040  
      id                 0.013680  
      Gender             0.012356  
      Gate location       0.000449  
      Unnamed: 0         -0.004552  
      Departure Delay in Minutes -0.050515  
      Departure/Arrival time convenient -0.051718  
      Arrival Delay in Minutes -0.057582  
      Customer Type      -0.187558  
      Type of Travel     -0.448995  
      Class              -0.449466  
      Name: satisfaction, dtype: float64
```

```
[35]: corr_matrix.style.background_gradient(cmap='coolwarm')
```

```
[35]: <pandas.io.formats.style.Styler at 0x11bbcc770>
```

1.1.3 Step 3: Split data into training set and test set

```
[39]: # Randomize the dataset  
      data_randomized = df.sample(frac=1, random_state=1)  
  
      # Calculate index for split - take first 80% of the data for test set  
      training_test_index = round(len(data_randomized) * 0.8)  
  
      # Split into training and test sets  
      training_set = data_randomized[:training_test_index].reset_index(drop=True)
```

```
test_set = data_randomized[training_test_index:].reset_index(drop=True)

print(training_set.shape)
print(test_set.shape)
```

(82875, 25)

(20719, 25)

Test to see if training set and test set proportions are approximately similar to original data set.

```
[42]: df['satisfaction'].value_counts(normalize=True)
```

```
[42]: satisfaction
0.0    0.566606
1.0    0.433394
Name: proportion, dtype: float64
```

```
[44]: training_set['satisfaction'].value_counts(normalize=True)
```

```
[44]: satisfaction
0.0    0.566287
1.0    0.433713
Name: proportion, dtype: float64
```

```
[46]: test_set['satisfaction'].value_counts(normalize=True)
```

```
[46]: satisfaction
0.0    0.567885
1.0    0.432115
Name: proportion, dtype: float64
```

1.1.4 Step 4: Apply machine learning methods and performance metrics

- Naive Bayes

```
[51]: trainX = training_set.iloc[:, :-1]
trainy = training_set['satisfaction']

testX = test_set.iloc[:, :-1]

testy = test_set['satisfaction']
NB_model = GaussianNB()
NB_model.fit(trainX, trainy)
y_pred = NB_model.predict(testX)

NB_accuracy = accuracy_score(testy, y_pred)
print(NB_accuracy)
```


0.8041893913798929

```
[53]: confusion_matrix(testy, y_pred)
```

```
[53]: array([[9564, 2202],  
         [1855, 7098]])
```

Even though the model has pretty good accuracy score, it seems the model is struggle with False Negative.

- **Logistic Regression**

```
[59]: from sklearn.linear_model import LogisticRegression  
log_reg_model = LogisticRegression(random_state=42)  
log_reg_model.fit(trainX, trainy)  
y_pred = log_reg_model.predict(testX)  
log_reg_accuracy = accuracy_score(testy, y_pred)  
print(log_reg_accuracy)
```

0.6843959650562286

```
/opt/anaconda3/lib/python3.12/site-  
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

- **Decision Tree**

```
[61]: from sklearn.tree import DecisionTreeClassifier  
dt_model = DecisionTreeClassifier(random_state=42, max_depth=5)  
dt_model.fit(trainX, trainy)  
y_pred = dt_model.predict(testX)  
dt_accuracy = accuracy_score(testy, y_pred)  
print(dt_accuracy)
```

0.90873111636662

It looks like I could improve the accuracy score by remove some less important feature columns. Based on correlation coefficients, I will drop the least correlated columns to “satisfaction” column: “id”, “Gender”, “Unnamed: 0”, “Gate location”, “Departure Delay in Minutes”, “Departure/Arrival time convenient”, “Arrival Delay in Minutes”

```
[67]: training_set = training_set.drop(columns=['id', 'Unnamed: 0', 'Gender', 'Gate_  
location', 'Departure Delay in Minutes',
```

```

'Departure/Arrival time convenient', 'Arrival Delay in_
↳Minutes'])

```

```

[69]: test_set = test_set.drop(columns=['id', 'Unnamed: 0', 'Gender', 'Gate_
↳location', 'Departure Delay in Minutes',
'Departure/Arrival time convenient', 'Arrival Delay in_
↳Minutes'])

```

I will perform all three models again to see if the accuracy score improves.

```

[84]: trainX = training_set.iloc[:, :-1]
trainy = training_set['satisfaction']
testX = test_set.iloc[:, :-1]
testy = test_set['satisfaction']
NB_model = GaussianNB()
NB_model.fit(trainX, trainy)
NB_y_pred = NB_model.predict(testX)
after_NB_accuracy = accuracy_score(testy, NB_y_pred)

log_reg_model = LogisticRegression(random_state=42)
log_reg_model.fit(trainX, trainy)
log_y_pred = log_reg_model.predict(testX)
after_log_reg_accuracy = accuracy_score(testy, log_y_pred)

dt_model = DecisionTreeClassifier(random_state=42, max_depth=5)
dt_model.fit(trainX, trainy)
dt_y_pred = dt_model.predict(testX)
after_dt_accuracy = accuracy_score(testy, dt_y_pred)

NB_cf = confusion_matrix(testy, NB_y_pred)
log_cf = confusion_matrix(testy, log_y_pred)
dt_cf = confusion_matrix(testy, dt_y_pred)

```

```

/opt/anaconda3/lib/python3.12/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(

```

```

[73]: print("Naive Bayes accuracy score:", NB_accuracy,
"\nLogistic Regression accuracy score", log_reg_accuracy,
"\nDecision Tree accuracy score", dt_accuracy,

```

```
"\n","\nNaive Bayes accuracy score AFTER:", after_NB_accuracy,
"\nLogistics Regression accuracy score AFTER:", after_log_reg_accuracy,
"\nDecision Tree accuracy score AFTER:", after_dt_accuracy)
```

Naive Bayes accuracy score: 0.8041893913798929

Logistic Regression accuracy score 0.6843959650562286

Decision Tree accuracy score 0.90873111636662

Naive Bayes accuracy score AFTER: 0.8651479318499927

Logistics Regression accuracy score AFTER: 0.8375886867126792

Decision Tree accuracy score AFTER: 0.908007143201892

```
[111]: NB_cf
```

```
[111]: array([[10672, 1094],
           [ 1700, 7253]])
```

```
[92]: log_cf
```

```
[92]: array([[10099, 1667],
           [ 1698, 7255]])
```

```
[98]: dt_cf
```

```
[98]: array([[10905, 861],
           [ 1045, 7908]])
```

Whether before or after remove less important feature variables, Decision Tree still performs the best out of 3 models. To further improve the performance of Decision Tree model, I will perform bagging to have a better performance for test set.

```
[79]: from sklearn.ensemble import BaggingClassifier
      bagg = BaggingClassifier(DecisionTreeClassifier(), n_estimators=500,
      ↪max_samples=100, n_jobs=-1, random_state=42)
      bagg.fit(trainX, trainy)
      y_pred = bagg.predict(testX)
      accuracy_score(testy, y_pred)
```

```
[79]: 0.9171292050774651
```

1.1.5 Conclusion

- High false negative or false positive in this case (airline industry) would not be as harmful as in medical field, where it could be dangerous and fatal. After removing less important features, the false positive is now higher than the false negative, even though the accuracy score is pretty high.

→ This could still greatly impact the airline business in many ways: not addressing the dissatisfaction of the customer in time could lead to the loss of loyalty, bad public reputation as customer might file complaint or address the issue to social media when they feel ignored.

- Specifically, Naïve Bayes model changes from false negative > false positive to false positive > false negative could indicate that even though the less important (less correlated) features to “Satisfaction level” are removed, the model might lack some context to distinguish between satisfy and dissatisfy classes.